

<  
(../03-  
basics-  
factors-  
dataframes/index.html)

Intro to R and RStudio for Genomics (../)

>  
(../05-  
data-  
visuali

## Aggregating and Analyzing Data with dplyr

### 🔗 Overview

**Teaching:** 40 min

**Exercises:** 15 min

#### Questions

- How can I manipulate dataframes without repeating myself?

#### Objectives

- Describe what the `dplyr` package in R is used for.
- Apply common `dplyr` functions to manipulate data in R.
- Employ the 'pipe' operator to link together a sequence of functions.
- Employ the 'mutate' function to apply other chosen functions to existing columns and create new columns of data.
- Employ the 'split-apply-combine' concept to split the data into groups, apply analysis to each group, and combine the results.

Bracket subsetting is handy, but it can be cumbersome and difficult to read, especially for complicated operations.

Luckily, the `[ dplyr ]` (<https://cran.r-project.org/package=dplyr>) package provides a number of very useful functions for manipulating dataframes in a way that will reduce repetition, reduce the probability of making errors, and probably even save you some typing. As an added bonus, you might even find the `dplyr` grammar easier to read.

Here we're going to cover 6 of the most commonly used functions as well as using pipes (`%>%`) to combine them.

1. `select()`
2. `filter()`
3. `group_by()`
4. `summarize()`
5. `mutate()`

Packages in R are sets of additional functions that let you do more stuff in R. The functions we've been using, like `str()`, come built into R; packages give you access to more functions. You need to install a package and then load it to be able to use it.

**R**

```
install.packages("dplyr") ## install
```

You might get asked to choose a CRAN mirror – this is asking you to choose a site to download the package from. The choice doesn't matter too much; I'd recommend choosing the RStudio mirror.

**R**

```
library("dplyr") ## Load
```

You only need to install a package once per computer, but you need to load it every time you open a new R session and want to use that package.

## What is dplyr?

The package `dplyr` is a fairly new (2014) package that tries to provide easy tools for the most common data manipulation tasks. It is built to work directly with data frames. The thinking behind it was largely inspired by the package `plyr` which has been in use for some time but suffered from being slow in some cases. `dplyr` addresses this by porting much of the computation to C++. An additional feature is the ability to work with data stored directly in an external database. The benefits of doing this are that the data can be managed natively in a relational database, queries can be conducted on that database, and only the results of the query returned.

This addresses a common problem with R in that all operations are conducted in memory and thus the amount of data you can work with is limited by available memory. The database connections essentially remove that limitation in that you can have a database of many 100s GB, conduct queries on it directly and pull back just what you need for analysis in R.

## Selecting columns and filtering rows

To select columns of a data frame, use `select()`. The first argument to this function is the data frame ( `variants` ), and the subsequent arguments are the columns to keep.

R

```
select(variants, sample_id, REF, ALT, DP)
```

### Output

```
# A tibble: 801 x 4
  sample_id REF      ALT      DP
  <chr>      <chr>    <chr>    <dbl>
1 SRR2584863 T        G         4
2 SRR2584863 G        T         6
3 SRR2584863 G        T        10
4 SRR2584863 CTTTTTTT CTTTTTTT  12
5 SRR2584863 CCGC     CCGCGC   10
6 SRR2584863 C        T        10
7 SRR2584863 C        A         8
8 SRR2584863 G        A        11
9 SRR2584863 ACAGCCAGCCAGCCAGCCA... ACAGCCAGCCAGCCAGCCAGCCAGCCAGC...  3
10 SRR2584863 AT       ATT        7
# ... with 791 more rows
```

To select all columns *except* certain ones, put a `"-"` in front of the variable to exclude it.

R

```
select(variants, -CHROM)
```

### Output

```
# A tibble: 801 x 28
  sample_id POS ID REF ALT QUAL FILTER INDEL IDV IMF DP
  <chr> <dbl> <lgl> <chr> <chr> <dbl> <lgl> <lgl> <dbl> <dbl> <dbl>
1 SRR25848... 9.97e3 NA T G 91 NA FALSE NA NA 4
2 SRR25848... 2.63e5 NA G T 85 NA FALSE NA NA 6
3 SRR25848... 2.82e5 NA G T 217 NA FALSE NA NA 10
4 SRR25848... 4.33e5 NA CTTT... CTTT... 64 NA TRUE 12 1 12
5 SRR25848... 4.74e5 NA CCGC CCGC... 228 NA TRUE 9 0.9 10
6 SRR25848... 6.49e5 NA C T 210 NA FALSE NA NA 10
7 SRR25848... 1.33e6 NA C A 178 NA FALSE NA NA 8
8 SRR25848... 1.73e6 NA G A 225 NA FALSE NA NA 11
9 SRR25848... 2.10e6 NA ACAG... ACAG... 56 NA TRUE 2 0.667 3
10 SRR25848... 2.33e6 NA AT ATT 167 NA TRUE 7 1 7
# ... with 791 more rows, and 17 more variables: VDB <dbl>, RPB <dbl>,
# MQB <dbl>, BQB <dbl>, MQSB <dbl>, SGB <dbl>, MQ0F <dbl>, ICB <lgl>,
# HOB <lgl>, AC <dbl>, AN <dbl>, DP4 <chr>, MQ <dbl>, Indiv <chr>,
# gt_PL <dbl>, gt_GT <dbl>, gt_GT_alleles <chr>
```

dplyr also provides useful functions to select columns based on their names. For instance, `ends_with()` allows you to select columns that ends with specific letters. For instance, if you wanted to select columns that end with the letter "B":

**R**

```
select(variants, ends_with("B"))
```

### Output

```
# A tibble: 801 x 8
  VDB RPB MQB BQB MQSB SGB ICB HOB
  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <lgl> <lgl>
1 0.0257 NA NA NA NA -0.556 NA NA
2 0.0961 1 1 1 NA -0.591 NA NA
3 0.774 NA NA NA 0.975 -0.662 NA NA
4 0.478 NA NA NA 1 -0.676 NA NA
5 0.660 NA NA NA 0.916 -0.662 NA NA
6 0.268 NA NA NA 0.916 -0.670 NA NA
7 0.624 NA NA NA 0.901 -0.651 NA NA
8 0.992 NA NA NA 1.01 -0.670 NA NA
9 0.902 NA NA NA 1 -0.454 NA NA
10 0.568 NA NA NA 1.01 -0.617 NA NA
# ... with 791 more rows
```

## Challenge

Create a table that contains all the columns with the letter “i” except for the columns “Indiv” and “FILTER”, and the column “POS”. Hint: look at the help for the function `ends_with()` we’ve just covered.

## Solution

R

```
select(variants, contains("i"), -Indiv, -FILTER, POS)
```

## Output

```
# A tibble: 801 x 7
  sample_id ID INDEL IDV IMF ICB POS
  <chr> <lg1> <lg1> <dbl> <dbl> <lg1> <dbl>
1 SRR2584863 NA FALSE NA NA NA 9972
2 SRR2584863 NA FALSE NA NA NA 263235
3 SRR2584863 NA FALSE NA NA NA 281923
4 SRR2584863 NA TRUE 12 1 NA 433359
5 SRR2584863 NA TRUE 9 0.9 NA 473901
6 SRR2584863 NA FALSE NA NA NA 648692
7 SRR2584863 NA FALSE NA NA NA 1331794
8 SRR2584863 NA FALSE NA NA NA 1733343
9 SRR2584863 NA TRUE 2 0.667 NA 2103887
10 SRR2584863 NA TRUE 7 1 NA 2333538
# ... with 791 more rows
```

To choose rows, we can use `filter()`. For instance, to keep the rows for the sample `SRR2584863`:

R

```
filter(variants, sample_id == "SRR2584863")
```

## Output

```
# A tibble: 25 x 29
  sample_id CHROM POS ID REF ALT QUAL FILTER INDEL IDV IMF
  <chr> <chr> <dbl> <lg1> <chr> <chr> <dbl> <lg1> <lg1> <dbl> <dbl>
1 SRR25848... CP00... 9.97e3 NA T G 91 NA FALSE NA NA
2 SRR25848... CP00... 2.63e5 NA G T 85 NA FALSE NA NA
3 SRR25848... CP00... 2.82e5 NA G T 217 NA FALSE NA NA
4 SRR25848... CP00... 4.33e5 NA CTTT... CTTT... 64 NA TRUE 12 1
5 SRR25848... CP00... 4.74e5 NA CCGC CCGC... 228 NA TRUE 9 0.9
6 SRR25848... CP00... 6.49e5 NA C T 210 NA FALSE NA NA
7 SRR25848... CP00... 1.33e6 NA C A 178 NA FALSE NA NA
8 SRR25848... CP00... 1.73e6 NA G A 225 NA FALSE NA NA
9 SRR25848... CP00... 2.10e6 NA ACAG... ACAG... 56 NA TRUE 2 0.667
10 SRR25848... CP00... 2.33e6 NA AT ATT 167 NA TRUE 7 1
# ... with 15 more rows, and 18 more variables: DP <dbl>, VDB <dbl>,
# RPB <dbl>, MQB <dbl>, BQB <dbl>, MQSB <dbl>, SGB <dbl>, MQOF <dbl>,
# ICB <lg1>, HOB <lg1>, AC <dbl>, AN <dbl>, DP4 <chr>, MQ <dbl>,
# Indiv <chr>, gt_PL <dbl>, gt_GT <dbl>, gt_GT_alleles <chr>
```

Note that this is equivalent to the base R code below, but is easier to read!

R

```
variants[variants$sample_id == "SRR2584863",]
```

`filter()` will keep all the rows that match the conditions that are provided. Here are a few examples:

R

```
## rows for which the reference genome has T or G
filter(variants, REF %in% c("T", "G"))
```

## Output

```
# A tibble: 340 x 29
  sample_id CHROM    POS ID   REF ALT   QUAL FILTER INDEL  IDV  IMF
  <chr>     <chr>  <dbl> <lgl> <chr> <chr> <dbl> <lgl> <lgl> <dbl> <dbl>
1 SRR25848... CP00... 9.97e3 NA   T    G     91   NA   FALSE NA    NA
2 SRR25848... CP00... 2.63e5 NA   G    T     85   NA   FALSE NA    NA
3 SRR25848... CP00... 2.82e5 NA   G    T    217   NA   FALSE NA    NA
4 SRR25848... CP00... 1.73e6 NA   G    A    225   NA   FALSE NA    NA
5 SRR25848... CP00... 2.62e6 NA   G    T    31.9 NA   FALSE NA    NA
6 SRR25848... CP00... 3.00e6 NA   G    A    225   NA   FALSE NA    NA
7 SRR25848... CP00... 3.91e6 NA   G    T    225   NA   FALSE NA    NA
8 SRR25848... CP00... 9.97e3 NA   T    G    214   NA   FALSE NA    NA
9 SRR25848... CP00... 1.00e4 NA   G    A    225   NA   FALSE NA    NA
10 SRR25848... CP00... 6.40e4 NA   G    A    225   NA   FALSE NA    NA
# ... with 330 more rows, and 18 more variables: DP <dbl>, VDB <dbl>,
#   RPB <dbl>, MQB <dbl>, BQB <dbl>, MQSB <dbl>, SGB <dbl>, MQ0F <dbl>,
#   ICB <lgl>, HOB <lgl>, AC <dbl>, AN <dbl>, DP4 <chr>, MQ <dbl>,
#   Indiv <chr>, gt_PL <dbl>, gt_GT <dbl>, gt_GT_alleles <chr>
```

R

```
## rows with QUAL values greater than or equal to 100
filter(variants, QUAL >= 100)
```

## Output

```
# A tibble: 666 x 29
  sample_id CHROM    POS ID   REF ALT   QUAL FILTER INDEL  IDV  IMF
  <chr>     <chr>  <dbl> <lgl> <chr> <chr> <dbl> <lgl> <lgl> <dbl> <dbl>
1 SRR25848... CP00... 2.82e5 NA   G    T    217   NA   FALSE NA    NA
2 SRR25848... CP00... 4.74e5 NA  CCGC CCGC... 228   NA   TRUE   9    0.9
3 SRR25848... CP00... 6.49e5 NA   C    T    210   NA   FALSE NA    NA
4 SRR25848... CP00... 1.33e6 NA   C    A    178   NA   FALSE NA    NA
5 SRR25848... CP00... 1.73e6 NA   G    A    225   NA   FALSE NA    NA
6 SRR25848... CP00... 2.33e6 NA  AT   ATT    167   NA   TRUE    7    1
7 SRR25848... CP00... 2.41e6 NA   A    C    104   NA   FALSE NA    NA
8 SRR25848... CP00... 2.45e6 NA   A    C    225   NA   FALSE NA    NA
9 SRR25848... CP00... 2.67e6 NA   A    T    225   NA   FALSE NA    NA
10 SRR25848... CP00... 3.00e6 NA   G    A    225   NA   FALSE NA    NA
# ... with 656 more rows, and 18 more variables: DP <dbl>, VDB <dbl>,
#   RPB <dbl>, MQB <dbl>, BQB <dbl>, MQSB <dbl>, SGB <dbl>, MQ0F <dbl>,
#   ICB <lgl>, HOB <lgl>, AC <dbl>, AN <dbl>, DP4 <chr>, MQ <dbl>,
#   Indiv <chr>, gt_PL <dbl>, gt_GT <dbl>, gt_GT_alleles <chr>
```

R

```
## rows that have TRUE in the column INDEL
filter(variants, INDEL)
```

## Output

```
# A tibble: 101 x 29
  sample_id CHROM   POS ID   REF  ALT   QUAL FILTER INDEL  IDV  IMF
  <chr>      <chr>   <dbl> <lgl> <chr> <chr> <dbl> <lgl> <lgl> <dbl> <dbl>
1 SRR25848... CP00... 4.33e5 NA   CTTT... CTTT... 64   NA   TRUE   12 1
2 SRR25848... CP00... 4.74e5 NA   CCGC  CCGC... 228  NA   TRUE    9 0.9
3 SRR25848... CP00... 2.10e6 NA   ACAG... ACAG... 56   NA   TRUE    2 0.667
4 SRR25848... CP00... 2.33e6 NA   AT    ATT   167  NA   TRUE    7 1
5 SRR25848... CP00... 3.90e6 NA   A     AC    43.4 NA   TRUE    2 1
6 SRR25848... CP00... 4.43e6 NA   TGG   T     228  NA   TRUE   10 1
7 SRR25848... CP00... 1.48e5 NA   AGGGG AGGG... 122  NA   TRUE    8 1
8 SRR25848... CP00... 1.58e5 NA   GTTT... GTTT... 19.5 NA   TRUE    6 1
9 SRR25848... CP00... 1.73e5 NA   CAA   CA    180  NA   TRUE   11 1
10 SRR25848... CP00... 1.75e5 NA   GAA   GA    194  NA   TRUE   10 1
# ... with 91 more rows, and 18 more variables: DP <dbl>, VDB <dbl>,
#   RPB <dbl>, MQB <dbl>, BQB <dbl>, MQSB <dbl>, SGB <dbl>, MQ0F <dbl>,
#   ICB <lgl>, HOB <lgl>, AC <dbl>, AN <dbl>, DP4 <chr>, MQ <dbl>,
#   Indiv <chr>, gt_PL <dbl>, gt_GT <dbl>, gt_GT_alleles <chr>
```

## R

```
## rows that don't have missing data in the IDV column
filter(variants, !is.na(IDV))
```

## Output

```
# A tibble: 101 x 29
  sample_id CHROM   POS ID   REF  ALT   QUAL FILTER INDEL  IDV  IMF
  <chr>      <chr>   <dbl> <lgl> <chr> <chr> <dbl> <lgl> <lgl> <dbl> <dbl>
1 SRR25848... CP00... 4.33e5 NA   CTTT... CTTT... 64   NA   TRUE   12 1
2 SRR25848... CP00... 4.74e5 NA   CCGC  CCGC... 228  NA   TRUE    9 0.9
3 SRR25848... CP00... 2.10e6 NA   ACAG... ACAG... 56   NA   TRUE    2 0.667
4 SRR25848... CP00... 2.33e6 NA   AT    ATT   167  NA   TRUE    7 1
5 SRR25848... CP00... 3.90e6 NA   A     AC    43.4 NA   TRUE    2 1
6 SRR25848... CP00... 4.43e6 NA   TGG   T     228  NA   TRUE   10 1
7 SRR25848... CP00... 1.48e5 NA   AGGGG AGGG... 122  NA   TRUE    8 1
8 SRR25848... CP00... 1.58e5 NA   GTTT... GTTT... 19.5 NA   TRUE    6 1
9 SRR25848... CP00... 1.73e5 NA   CAA   CA    180  NA   TRUE   11 1
10 SRR25848... CP00... 1.75e5 NA   GAA   GA    194  NA   TRUE   10 1
# ... with 91 more rows, and 18 more variables: DP <dbl>, VDB <dbl>,
#   RPB <dbl>, MQB <dbl>, BQB <dbl>, MQSB <dbl>, SGB <dbl>, MQ0F <dbl>,
#   ICB <lgl>, HOB <lgl>, AC <dbl>, AN <dbl>, DP4 <chr>, MQ <dbl>,
#   Indiv <chr>, gt_PL <dbl>, gt_GT <dbl>, gt_GT_alleles <chr>
```

`filter()` allows you to combine multiple conditions. You can separate them using a `,` as arguments to the function, they will be combined using the `&` (AND) logical operator. If you need to use the `|` (OR) logical operator, you can specify it explicitly:

## R

```
## this is equivalent to:
## filter(variants, sample_id == "SRR2584863" & QUAL >= 100)
filter(variants, sample_id == "SRR2584863", QUAL >= 100)
```

## Output

```
# A tibble: 19 x 29
  sample_id CHROM    POS ID   REF  ALT   QUAL FILTER INDEL  IDV  IMF
  <chr>      <chr>    <dbl> <lgl> <chr> <chr> <dbl> <lgl> <lgl> <dbl> <dbl>
1 SRR25848... CP00... 2.82e5 NA    G    T     217 NA    FALSE NA    NA
2 SRR25848... CP00... 4.74e5 NA    CCGC CCGC... 228 NA    TRUE   9    0.9
3 SRR25848... CP00... 6.49e5 NA    C    T     210 NA    FALSE NA    NA
4 SRR25848... CP00... 1.33e6 NA    C    A     178 NA    FALSE NA    NA
5 SRR25848... CP00... 1.73e6 NA    G    A     225 NA    FALSE NA    NA
6 SRR25848... CP00... 2.33e6 NA    AT   ATT    167 NA    TRUE   7    1
7 SRR25848... CP00... 2.41e6 NA    A    C     104 NA    FALSE NA    NA
8 SRR25848... CP00... 2.45e6 NA    A    C     225 NA    FALSE NA    NA
9 SRR25848... CP00... 2.67e6 NA    A    T     225 NA    FALSE NA    NA
10 SRR25848... CP00... 3.00e6 NA    G    A     225 NA    FALSE NA    NA
11 SRR25848... CP00... 3.34e6 NA    A    C     211 NA    FALSE NA    NA
12 SRR25848... CP00... 3.40e6 NA    C    A     225 NA    FALSE NA    NA
13 SRR25848... CP00... 3.48e6 NA    A    G     200 NA    FALSE NA    NA
14 SRR25848... CP00... 3.49e6 NA    A    C     225 NA    FALSE NA    NA
15 SRR25848... CP00... 3.91e6 NA    G    T     225 NA    FALSE NA    NA
16 SRR25848... CP00... 4.10e6 NA    A    G     225 NA    FALSE NA    NA
17 SRR25848... CP00... 4.20e6 NA    A    C     225 NA    FALSE NA    NA
18 SRR25848... CP00... 4.43e6 NA    TGG   T     228 NA    TRUE  10    1
19 SRR25848... CP00... 4.62e6 NA    A    C     185 NA    FALSE NA    NA
# ... with 18 more variables: DP <dbl>, VDB <dbl>, RPB <dbl>, MQB <dbl>,
#   BQB <dbl>, MQSB <dbl>, SGB <dbl>, MQ0F <dbl>, ICB <lgl>, HOB <lgl>,
#   AC <dbl>, AN <dbl>, DP4 <chr>, MQ <dbl>, Indiv <chr>, gt_PL <dbl>,
#   gt_GT <dbl>, gt_GT_alleles <chr>
```

## R

```
## using `|` Logical operator
filter(variants, sample_id == "SRR2584863", (INDEL | QUAL >= 100))
```

## Output

```
# A tibble: 22 x 29
  sample_id CHROM    POS ID   REF  ALT   QUAL FILTER INDEL  IDV  IMF
  <chr>      <chr>    <dbl> <lgl> <chr> <chr> <dbl> <lgl> <lgl> <dbl> <dbl>
1 SRR25848... CP00... 2.82e5 NA    G    T     217 NA    FALSE NA    NA
2 SRR25848... CP00... 4.33e5 NA    CTTT... CTTT... 64 NA    TRUE  12    1
3 SRR25848... CP00... 4.74e5 NA    CCGC CCGC... 228 NA    TRUE   9    0.9
4 SRR25848... CP00... 6.49e5 NA    C    T     210 NA    FALSE NA    NA
5 SRR25848... CP00... 1.33e6 NA    C    A     178 NA    FALSE NA    NA
6 SRR25848... CP00... 1.73e6 NA    G    A     225 NA    FALSE NA    NA
7 SRR25848... CP00... 2.10e6 NA    ACAG... ACAG... 56 NA    TRUE   2    0.667
8 SRR25848... CP00... 2.33e6 NA    AT   ATT    167 NA    TRUE   7    1
9 SRR25848... CP00... 2.41e6 NA    A    C     104 NA    FALSE NA    NA
10 SRR25848... CP00... 2.45e6 NA    A    C     225 NA    FALSE NA    NA
# ... with 12 more rows, and 18 more variables: DP <dbl>, VDB <dbl>,
#   RPB <dbl>, MQB <dbl>, BQB <dbl>, MQSB <dbl>, SGB <dbl>, MQ0F <dbl>,
#   ICB <lgl>, HOB <lgl>, AC <dbl>, AN <dbl>, DP4 <chr>, MQ <dbl>,
#   Indiv <chr>, gt_PL <dbl>, gt_GT <dbl>, gt_GT_alleles <chr>
```

## Challenge

Select all the mutations that occurred between the positions 1e6 (one million) and 2e6 (included) that are not indels and have QUAL greater than 200.

## Solution

### R

```
filter(variants, POS >= 1e6 & POS <= 2e6, !INDEL, QUAL > 200)
```

### Output

```
# A tibble: 77 x 29
  sample_id CHROM  POS ID  REF  ALT  QUAL FILTER INDEL  IDV  IMF
<chr>      <chr> <dbl> <lgl> <chr> <chr> <dbl> <lgl> <lgl> <dbl> <dbl>
1 SRR25848... CP00... 1.73e6 NA  G   A    225 NA   FALSE NA   NA
2 SRR25848... CP00... 1.00e6 NA  A   G    225 NA   FALSE NA   NA
3 SRR25848... CP00... 1.02e6 NA  A   G    225 NA   FALSE NA   NA
4 SRR25848... CP00... 1.06e6 NA  C   T    225 NA   FALSE NA   NA
5 SRR25848... CP00... 1.06e6 NA  A   G    206 NA   FALSE NA   NA
6 SRR25848... CP00... 1.07e6 NA  G   T    225 NA   FALSE NA   NA
7 SRR25848... CP00... 1.07e6 NA  T   C    225 NA   FALSE NA   NA
8 SRR25848... CP00... 1.10e6 NA  C   T    225 NA   FALSE NA   NA
9 SRR25848... CP00... 1.11e6 NA  C   T    212 NA   FALSE NA   NA
10 SRR25848... CP00... 1.11e6 NA  A   G    225 NA   FALSE NA   NA
# ... with 67 more rows, and 18 more variables: DP <dbl>, VDB <dbl>,
# RPB <dbl>, MQB <dbl>, BQB <dbl>, MQSB <dbl>, SGB <dbl>, MQ0F <dbl>,
# ICB <lgl>, HOB <lgl>, AC <dbl>, AN <dbl>, DP4 <chr>, MQ <dbl>,
# Indiv <chr>, gt_PL <dbl>, gt_GT <dbl>, gt_GT_alleles <chr>
```

## Pipes

But what if you wanted to select and filter? We can do this with pipes. Pipes, are a fairly recent addition to R. Pipes let you take the output of one function and send it directly to the next, which is useful when you need to many things to the same data set. It was possible to do this before pipes were added to R, but it was much messier and more difficult. Pipes in R look like `%>%` and are made available via the `magrittr` package, which is installed as part of `dplyr`. If you use RStudio, you can type the pipe with `Ctrl + Shift + M` if you're using a PC, or `Cmd + Shift + M` if you're using a Mac.

### R

```
variants %>%
  filter(sample_id == "SRR2584863") %>%
  select(REF, ALT, DP)
```

### Output



```
# A tibble: 25 x 3
  REF      ALT      DP
  <chr>   <chr>   <dbl>
1 T       G         4
2 G       T         6
3 G       T        10
4 CTTTCTTT CTTTCTTT  12
5 CCGC    CCGCGC   10
6 C       T        10
7 C       A         8
8 G       A        11
9 ACAGCCAGCCAGCCAGCCAGCCAG... ACAGCCAGCCAGCCAGCCAGCCAGCCAGCCAGCCA... 3
10 AT     ATT         7
# ... with 15 more rows
```

In the above code, we use the pipe to send the `variants` dataset first through `filter()`, to keep rows where `sample_id` matches a particular sample, and then through `select()` to keep only the `REF`, `ALT`, and `DP` columns. Since `%>%` takes the object on its left and passes it as the first argument to the function on its right, we don't need to explicitly include the data frame as an argument to the `filter()` and `select()` functions any more. We then pipe the results to the `head()` function so that we only see the first six rows of data.

Some may find it helpful to read the pipe like the word "then". For instance, in the above example, we took the data frame `variants`, *then* we filtered for rows where `sample_id` was `SRR2584863`, *then* we selected the `REF`, `ALT`, and `DP` columns, *then* we showed only the first six rows. The `dplyr` functions by themselves are somewhat simple, but by combining them into linear workflows with the pipe, we can accomplish more complex manipulations of data frames.

If we want to create a new object with this smaller version of the data we can do so by assigning it a new name:

**R**

```
SRR2584863_variants <- variants %>%
  filter(sample_id == "SRR2584863") %>%
  select(REF, ALT, DP)
```

This new object includes all of the data from this sample. Let's look at it to confirm it's what we want:

**R**

```
SRR2584863_variants
```

### Output

```
# A tibble: 25 x 3
  REF      ALT      DP
  <chr>   <chr>   <dbl>
1 T       G         4
2 G       T         6
3 G       T        10
4 CTTTCTTT CTTTCTTT  12
5 CCGC    CCGCGC   10
6 C       T        10
7 C       A         8
8 G       A        11
9 ACAGCCAGCCAGCCAGCCAGCCAG... ACAGCCAGCCAGCCAGCCAGCCAGCCAGCCAGCCA... 3
10 AT     ATT         7
# ... with 15 more rows
```

## Exercise: Pipe and filter

Starting with the `variants` dataframe, use pipes to subset the data to include only observations from SRR2584863 sample, where the filtered depth (DP) is at least 10. Retain only the columns `REF`, `ALT`, and `POS`.

### Solution

#### R

```
variants %>%
  filter(sample_id == "SRR2584863" & DP >= 10) %>%
  select(REF, ALT, POS)
```

#### Output

```
# A tibble: 16 x 3
  REF      ALT      POS
  <chr>   <chr>   <dbl>
1 G       T       281923
2 CTTTCTT CTTTCTT  433359
3 CCGC    CCGCGC   473901
4 C       T       648692
5 G       A       1733343
6 A       C       2446984
7 G       T       2618472
8 A       T       2665639
9 G       A       2999330
10 A      C       3339313
11 C      A       3401754
12 A      C       3488669
13 G      T       3909807
14 A      G       4100183
15 A      C       4201958
16 TGG    T       4431393
```

## Mutate

Frequently you'll want to create new columns based on the values in existing columns, for example to do unit conversions or find the ratio of values in two columns. For this we'll use the `dplyr` function `mutate()`.

We have a column titled "QUAL". This is a Phred-scaled confidence score that a polymorphism exists at this position given the sequencing data. Lower QUAL scores indicate low probability of a polymorphism existing at that site. We can convert the confidence value QUAL to a probability value according to the formula:

$$\text{Probability} = 1 - 10^{-(\text{QUAL}/10)}$$

Let's add a column (`POLPROB`) to our `variants` dataframe that shows the probability of a polymorphism at that site given the data. We'll show only the first six rows of data.

#### R

```
variants %>%
  mutate(POLPROB = 1 - (10 ^ -(QUAL/10)))
```

#### Output

```
# A tibble: 801 x 30
  sample_id CHROM    POS ID   REF  ALT   QUAL FILTER INDEL  IDV   IMF
  <chr>      <chr>  <dbl> <lgl> <chr> <chr> <dbl> <lgl> <lgl> <dbl> <dbl>
1 SRR25848... CP00... 9.97e3 NA   T    G     91 NA   FALSE NA NA
2 SRR25848... CP00... 2.63e5 NA   G    T     85 NA   FALSE NA NA
3 SRR25848... CP00... 2.82e5 NA   G    T    217 NA   FALSE NA NA
4 SRR25848... CP00... 4.33e5 NA   CTTT... CTTT... 64 NA   TRUE  12  1
5 SRR25848... CP00... 4.74e5 NA   CCGC  CCGC... 228 NA   TRUE   9  0.9
6 SRR25848... CP00... 6.49e5 NA   C    T    210 NA   FALSE NA NA
7 SRR25848... CP00... 1.33e6 NA   C    A    178 NA   FALSE NA NA
8 SRR25848... CP00... 1.73e6 NA   G    A    225 NA   FALSE NA NA
9 SRR25848... CP00... 2.10e6 NA   ACAG... ACAG... 56 NA   TRUE   2  0.667
10 SRR25848... CP00... 2.33e6 NA   AT    ATT   167 NA   TRUE   7  1
# ... with 791 more rows, and 19 more variables: DP <dbl>, VDB <dbl>,
# RPB <dbl>, MQB <dbl>, BQB <dbl>, MQSB <dbl>, SGB <dbl>, MQ0F <dbl>,
# ICB <lgl>, HOB <lgl>, AC <dbl>, AN <dbl>, DP4 <chr>, MQ <dbl>,
# Indiv <chr>, gt_PL <dbl>, gt_GT <dbl>, gt_GT_alleles <chr>,
# POLPROB <dbl>
```

## Exercise

There are a lot of columns in our dataset, so let's just look at the `sample_id`, `POS`, `QUAL`, and `POLPROB` columns for now. Add a line to the above code to only show those columns.

## Solution

### R

```
variants %>%
  mutate(POLPROB = 1 - 10 ^ -(QUAL/10)) %>%
  select(sample_id, POS, QUAL, POLPROB)
```

### Output

```
# A tibble: 801 x 4
  sample_id    POS  QUAL POLPROB
  <chr>      <dbl> <dbl> <dbl>
1 SRR2584863  9972   91  1.000
2 SRR2584863 263235   85  1.000
3 SRR2584863 281923  217    1
4 SRR2584863 433359   64  1.000
5 SRR2584863 473901  228    1
6 SRR2584863 648692  210    1
7 SRR2584863 1331794 178    1
8 SRR2584863 1733343 225    1
9 SRR2584863 2103887   56  1.000
10 SRR2584863 2333538 167    1
# ... with 791 more rows
```

We are interested in knowing the most common size for the indels. Let's create a new column, called "indel\_size" that contains the size difference between the our sequences and the reference genome. The function, `nchar()` returns the number of letters in a string.

### R

```
variants %>%
  mutate(indel_size = nchar(ALT) - nchar(REF))
```

### Output

```
# A tibble: 801 x 30
  sample_id CHROM   POS ID   REF  ALT   QUAL FILTER INDEL  IDV   IMF
  <chr>      <chr>   <dbl> <lgl> <chr> <chr> <dbl> <lgl> <lgl> <dbl> <dbl>
1 SRR25848... CP00... 9.97e3 NA   T    G     91 NA   FALSE NA NA
2 SRR25848... CP00... 2.63e5 NA   G    T     85 NA   FALSE NA NA
3 SRR25848... CP00... 2.82e5 NA   G    T    217 NA   FALSE NA NA
4 SRR25848... CP00... 4.33e5 NA   CTTT... CTTT... 64 NA   TRUE  12  1
5 SRR25848... CP00... 4.74e5 NA   CCGC CCGC... 228 NA   TRUE   9  0.9
6 SRR25848... CP00... 6.49e5 NA   C    T    210 NA   FALSE NA NA
7 SRR25848... CP00... 1.33e6 NA   C    A    178 NA   FALSE NA NA
8 SRR25848... CP00... 1.73e6 NA   G    A    225 NA   FALSE NA NA
9 SRR25848... CP00... 2.10e6 NA   ACAG... ACAG... 56 NA   TRUE   2  0.667
10 SRR25848... CP00... 2.33e6 NA   AT    ATT  167 NA   TRUE   7  1
# ... with 791 more rows, and 19 more variables: DP <dbl>, VDB <dbl>,
# RPB <dbl>, MQB <dbl>, BQB <dbl>, MQSB <dbl>, SGB <dbl>, MQ0F <dbl>,
# ICB <lgl>, HOB <lgl>, AC <dbl>, AN <dbl>, DP4 <chr>, MQ <dbl>,
# Indiv <chr>, gt_PL <dbl>, gt_GT <dbl>, gt_GT_alleles <chr>,
# indel_size <int>
```

When you want to create a new variable that depends on multiple conditions, the function `case_when()` in combination with `mutate()` is very useful. Our current dataset has a column that tells us whether each mutation is an indel, but we don't know if it's an insertion or a deletion. Let's create a new variable, called `mutation_type` that will take the values: `insertion`, `deletion`, or `point` depending on the value found in the `indel_size` column. We will save this data frame in a new variable, called `variants_indel`.

R

```
variants_indel <- variants %>%
  mutate(
    indel_size = nchar(ALT) - nchar(REF),
    mutation_type = case_when(
      indel_size > 0 ~ "insertion",
      indel_size < 0 ~ "deletion",
      indel_size == 0 ~ "point"
    )
  )
```

When `case_when()` is used within `mutate()`, each row is evaluated for the condition listed in the order listed. The first condition that returns `TRUE` will be used to fill the content of the new column (here `mutation_type`) with the value listed on the right side of the `~` is used. If none of the conditions are met, the function returns `NA` (missing data).

We can check that we captured all possibilities by looking for missing data in the new `mutation_type` column, and confirm that no row matches this condition:

R

```
variants_indel %>%
  filter(is.na(mutation_type))
```

Output

```
# A tibble: 0 x 31
# ... with 31 variables: sample_id <chr>, CHROM <chr>, POS <dbl>, ID <lgl>,
# REF <chr>, ALT <chr>, QUAL <dbl>, FILTER <lgl>, INDEL <lgl>,
# IDV <dbl>, IMF <dbl>, DP <dbl>, VDB <dbl>, RPB <dbl>, MQB <dbl>,
# BQB <dbl>, MQSB <dbl>, SGB <dbl>, MQ0F <dbl>, ICB <lgl>, HOB <lgl>,
# AC <dbl>, AN <dbl>, DP4 <chr>, MQ <dbl>, Indiv <chr>, gt_PL <dbl>,
# gt_GT <dbl>, gt_GT_alleles <chr>, indel_size <int>,
# mutation_type <chr>
```

## Split-apply-combine data analysis and the summarize() function

Many data analysis tasks can be approached using the “split-apply-combine” paradigm: split the data into groups, apply some analysis to each group, and then combine the results. `dplyr` makes this very easy through the use of the `group_by()` function, which splits the data into groups. When the data is grouped in this way `summarize()` can be used to collapse each group into a single-row summary. `summarize()` does this by applying an aggregating or summary function to each group.

For example, if we wanted to group by `mutation_type` and find the average size for the insertions and deletions:

**R**

```
variants_indel %>%
  group_by(mutation_type) %>%
  summarize(
    mean_size = mean(indel_size)
  )
```

**Output**

```
# A tibble: 3 x 2
  mutation_type mean_size
  <chr>         <dbl>
1 deletion      -1.38
2 insertion      1.61
3 point          0
```

We can have additional columns by adding arguments to the `summarize()` function. For instance, if we also wanted to know the median indel size:

**R**

```
variants_indel %>%
  group_by(mutation_type) %>%
  summarize(
    mean_size = mean(indel_size),
    median_size = median(indel_size)
  )
```

**Output**

```
# A tibble: 3 x 3
  mutation_type mean_size median_size
  <chr>         <dbl>         <dbl>
1 deletion      -1.38          -1
2 insertion      1.61           1
3 point          0           0
```

So to view the highest filtered depth ( `DP` ) for each sample:

**R**

```
variants_indel %>%
  group_by(sample_id) %>%
  summarize(max(DP))
```

**Output**

```
# A tibble: 3 x 2
  sample_id `max(DP)`
  <chr>         <dbl>
1 SRR2584863    20
2 SRR2584866    79
3 SRR2589044    16
```

## Challenge

What are the largest insertions and deletions? Hint: the function `abs()` returns the absolute value.

## Solution

### R

```
variants_indel %>%
  group_by(mutation_type) %>%
  summarize(
    max_size = max(abs(indel_size))
  )
```

### Output

```
# A tibble: 3 x 2
  mutation_type max_size
  <chr>         <int>
1 deletion           6
2 insertion          24
3 point              0
```

## Callout: missing data and built-in functions

R has many built-in functions like `mean()`, `median()`, `min()`, and `max()` that are useful to compute summary statistics. These are called “built-in functions” because they come with R and don’t require that you install any additional packages. By default, all **R functions operating on vectors that contains missing data will return NA**. It’s a way to make sure that users know they have missing data, and make a conscious decision on how to deal with it. When dealing with simple statistics like the mean, the easiest way to ignore NA (the missing data) is to use `na.rm = TRUE` (`rm` stands for remove).

It is often useful to calculate how many observations are present in each group. The function `n()` helps you do that:

### R

```
variants_indel %>%
  group_by(mutation_type) %>%
  summarize(
    n = n()
  )
```

### Output

```
# A tibble: 3 x 2
  mutation_type    n
  <chr>         <int>
1 deletion        39
2 insertion        62
3 point          700
```

Because it’s a common operation, the `dplyr` verb, `count()` is a “shortcut” that combines these 2 commands:

### R

```
variants_indel %>%
  count(mutation_type)
```

### Output

```
# A tibble: 3 x 2
  mutation_type      n
  <chr>          <int>
1 deletion         39
2 insertion         62
3 point           700
```

This lesson is in the early stages of development (Alpha version)

`group_by()` (and therefore `count()`) can also take multiple column names.

## Challenge

- How many mutations are found in each sample?

### Solution

**R**

```
variants_indel %>%
  count(sample_id)
```

#### Output

```
# A tibble: 3 x 2
  sample_id      n
  <chr>        <int>
1 SRR2584863    25
2 SRR2584866   766
3 SRR2589044    10
```

- How many mutation types are found in sample?

### Solution

**R**

```
variants_indel %>%
  count(sample_id, mutation_type)
```

#### Output

```
# A tibble: 9 x 3
  sample_id mutation_type      n
  <chr>      <chr>        <int>
1 SRR2584863 deletion          1
2 SRR2584863 insertion          5
3 SRR2584863 point           19
4 SRR2584866 deletion         37
5 SRR2584866 insertion        54
6 SRR2584866 point          675
7 SRR2589044 deletion          1
8 SRR2589044 insertion          3
9 SRR2589044 point            6
```

## Reshaping data frames

While the tidy format is useful to analyze and plot data in R, it can sometimes be useful to transform the “long” tidy format, into the wide format. This transformation can be done with the `spread()` function provided by the `tidyr` package (also part of the `tidyverse` ).

`spread()` takes a data frame as the first argument, and two arguments: the column name that will become the columns and the column name that will become the cells in the wide data.

**R**

```
variants_wide <- variants_indel %>%  
  count(sample_id, mutation_type) %>%  
  spread(mutation_type, n)  
variants_wide
```

**Output**

```
# A tibble: 3 x 4  
  sample_id deletion insertion point  
  <chr>      <int>    <int> <int>  
1 SRR2584863      1        5    19  
2 SRR2584866     37       54   675  
3 SRR2589044      1        3     6
```

The opposite operation of `spread()` is taken care by `gather()` . We specify the names of the new columns, and here add `-sample_id` as this column shouldn't be affected by the reshaping:

**R**

```
variants_wide %>%  
  gather(mutation_type, n, -sample_id)
```

**Output**

```
# A tibble: 9 x 3  
  sample_id mutation_type      n  
  <chr>      <chr>    <int>  
1 SRR2584863 deletion        1  
2 SRR2584866 deletion       37  
3 SRR2589044 deletion        1  
4 SRR2584863 insertion        5  
5 SRR2584866 insertion       54  
6 SRR2589044 insertion        3  
7 SRR2584863 point         19  
8 SRR2584866 point        675  
9 SRR2589044 point          6
```



## ✎ Challenge (optional)

Based on the probability scores we calculated when we first introduced `mutate()`, classify each mutation in 3 categories: high ( $> 0.95$ ), medium (between 0.7 and 0.95), and low ( $< 0.7$ ), and create a table with `sample_id` as rows, the 3 levels of quality as columns, and the number of mutations in the cells.

## 👁 Solution

### R

```
variants %>%
  mutate(POLPROB = 1 - (10 ^ -(QUAL/10))) %>%
  mutate(prob_cat = case_when(
    POLPROB >= .95 ~ "high",
    POLPROB >= .7 & POLPROB < .95 ~ "medium",
    POLPROB < .7 ~ "low"
  )) %>%
  count(sample_id, prob_cat) %>%
  spread(prob_cat, n)
```

### Output

```
# A tibble: 3 x 4
  sample_id  high  low medium
  <chr>      <int> <int> <int>
1 SRR2584863    25   NA     NA
2 SRR2584866   754    4      8
3 SRR2589044    10   NA     NA
```

## Exporting

We can export this new dataset using `write_csv()` :

### R

```
write_csv(variants_indel, "variants_indel.csv")
```

## Resources

- Handy dplyr cheatsheet (<https://github.com/rstudio/cheatsheets/raw/master/data-transformation.pdf>)

*Much of this lesson was copied or adapted from Jeff Hollister's materials (<http://usepa.github.io/introR/2015/01/14/03-Clean/>)*

## 📌 Key Points

- Use the `dplyr` package to manipulate dataframes.
- Use `select()` to choose variables from a dataframe.
- Use `filter()` to choose data based on values.
- Use `group_by()` and `summarize()` to work with subsets of data.
- Use `mutate()` to create new variables.

<  
(../03-  
basics-  
factors-  
dataframes/index.html)

>  
(../05-  
data-  
visuali

---

Licensed under CC-BY 4.0 (<https://creativecommons.org/licenses/by/4.0/>) 2018–2019 by The Carpentries (<https://carpentries.org/>)  
Licensed under CC-BY 4.0 (<https://creativecommons.org/licenses/by/4.0/>) 2016–2018 by Data Carpentry (<http://datacarpentry.org>)

Edit on GitHub ([https://github.com/datacarpentry/genomics-r-intro/edit/master/\\_episodes\\_rmd/04-dplyr.Rmd](https://github.com/datacarpentry/genomics-r-intro/edit/master/_episodes_rmd/04-dplyr.Rmd)) / Contributing  
(<https://github.com/datacarpentry/genomics-r-intro/blob/gh-pages/CONTRIBUTING.md>) / Source  
(<https://github.com/datacarpentry/genomics-r-intro/>) / Cite (<https://github.com/datacarpentry/genomics-r-intro/blob/gh-pages/CITATION>) /  
Contact (<mailto:team@carpentries.org>)

Using The Carpentries theme (<https://github.com/carpentries/carpentries-theme/>) — Site last built on: 2019-06-02 01:03:53 +0000.